

Multimodal Document Intelligence Assistant

This project is a multimodal, multi-document Retrieval-Augmented Generation (RAG) system designed to provide accurate, grounded answers exclusively from uploaded documents. Unlike traditional PDF chatbots, this system handles layout structures, tables, and images by leveraging hybrid database querying and vision-language model descriptions.

Key System Performance Metrics

Below is a table showing the performance benchmarks of the RAG tool components:

Component	Latency (s)	Accuracy (%)
Text Chunk Search	0.15	98.5
Table Understanding	0.45	94.2
Image Captioning	1.80	89.0